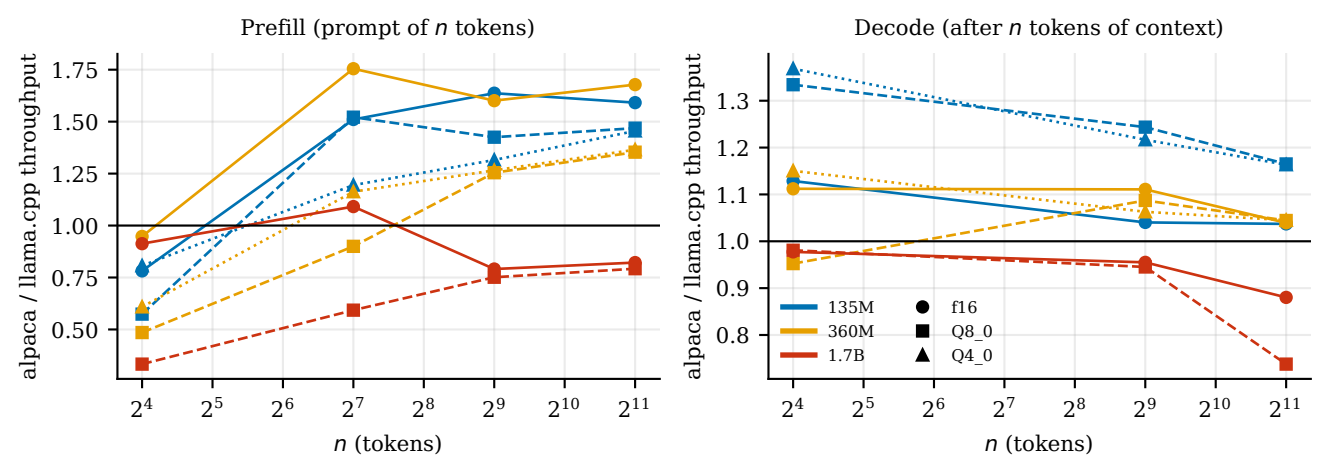
Prefill (prompt of n tokens)
alpaca / llama.cpp throughput
1.75
1.50
1.25
1.00
0.75
0.50
2^4 2^5 2^6 2^7 2^8 2^9 2^{10} 2^{11}
n (tokens)
Decode (after n tokens of context)
alpaca / llama.cpp throughput
1.3
1.2
1.1
1.0
0.9
0.8
135M
360M
1.7B
f16
Q8_0
Q4_0
2^4 2^5 2^6 2^7 2^8 2^9 2^{10} 2^{11}
n (tokens)